

HousePrediction

April 23, 2019

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: d=pd.read_csv(r'D:\HousePredictions\train.csv',encoding='unicode_escape')
d.head()
```

```
Out[2]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000

[5 rows x 81 columns]

```
In [3]: d.describe()
```

```
Out[3]:
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	\
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	
std	421.610009	42.300571	24.284752	9981.264932	1.382997	
min	1.000000	20.000000	21.000000	1300.000000	1.000000	
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	

75%	1095.250000	70.000000	80.000000	11601.500000	7.000000
max	1460.000000	190.000000	313.000000	215245.000000	10.000000

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	\
count	1460.000000	1460.000000	1460.000000	1452.000000	1460.000000	
mean	5.575342	1971.267808	1984.865753	103.685262	443.639726	
std	1.112799	30.202904	20.645407	181.066207	456.098091	
min	1.000000	1872.000000	1950.000000	0.000000	0.000000	
25%	5.000000	1954.000000	1967.000000	0.000000	0.000000	
50%	5.000000	1973.000000	1994.000000	0.000000	383.500000	
75%	6.000000	2000.000000	2004.000000	166.000000	712.250000	
max	9.000000	2010.000000	2010.000000	1600.000000	5644.000000	

	...	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch	\
count	...	1460.000000	1460.000000	1460.000000	1460.000000	
mean	...	94.244521	46.660274	21.954110	3.409589	
std	...	125.338794	66.256028	61.119149	29.317331	
min	...	0.000000	0.000000	0.000000	0.000000	
25%	...	0.000000	0.000000	0.000000	0.000000	
50%	...	0.000000	25.000000	0.000000	0.000000	
75%	...	168.000000	68.000000	0.000000	0.000000	
max	...	857.000000	547.000000	552.000000	508.000000	

	ScreenPorch	PoolArea	MiscVal	MoSold	YrSold	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	15.060959	2.758904	43.489041	6.321918	2007.815753	
std	55.757415	40.177307	496.123024	2.703626	1.328095	
min	0.000000	0.000000	0.000000	1.000000	2006.000000	
25%	0.000000	0.000000	0.000000	5.000000	2007.000000	
50%	0.000000	0.000000	0.000000	6.000000	2008.000000	
75%	0.000000	0.000000	0.000000	8.000000	2009.000000	
max	480.000000	738.000000	15500.000000	12.000000	2010.000000	

	SalePrice
count	1460.000000
mean	180921.195890
std	79442.502883
min	34900.000000
25%	129975.000000
50%	163000.000000
75%	214000.000000
max	755000.000000

[8 rows x 38 columns]

In [4]: d.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
```

Data columns (total 81 columns):

Id	1460 non-null int64
MSSubClass	1460 non-null int64
MSZoning	1460 non-null object
LotFrontage	1201 non-null float64
LotArea	1460 non-null int64
Street	1460 non-null object
Alley	91 non-null object
LotShape	1460 non-null object
LandContour	1460 non-null object
Utilities	1460 non-null object
LotConfig	1460 non-null object
LandSlope	1460 non-null object
Neighborhood	1460 non-null object
Condition1	1460 non-null object
Condition2	1460 non-null object
BldgType	1460 non-null object
HouseStyle	1460 non-null object
OverallQual	1460 non-null int64
OverallCond	1460 non-null int64
YearBuilt	1460 non-null int64
YearRemodAdd	1460 non-null int64
RoofStyle	1460 non-null object
RoofMatl	1460 non-null object
Exterior1st	1460 non-null object
Exterior2nd	1460 non-null object
MasVnrType	1452 non-null object
MasVnrArea	1452 non-null float64
ExterQual	1460 non-null object
ExterCond	1460 non-null object
Foundation	1460 non-null object
BsmtQual	1423 non-null object
BsmtCond	1423 non-null object
BsmtExposure	1422 non-null object
BsmtFinType1	1423 non-null object
BsmtFinSF1	1460 non-null int64
BsmtFinType2	1422 non-null object
BsmtFinSF2	1460 non-null int64
BsmtUnfSF	1460 non-null int64
TotalBsmtSF	1460 non-null int64
Heating	1460 non-null object
HeatingQC	1460 non-null object
CentralAir	1460 non-null object
Electrical	1459 non-null object
1stFlrSF	1460 non-null int64
2ndFlrSF	1460 non-null int64
LowQualFinSF	1460 non-null int64
GrLivArea	1460 non-null int64

```

BsmtFullBath      1460 non-null int64
BsmtHalfBath      1460 non-null int64
FullBath          1460 non-null int64
HalfBath          1460 non-null int64
BedroomAbvGr      1460 non-null int64
KitchenAbvGr      1460 non-null int64
KitchenQual       1460 non-null object
TotRmsAbvGrd      1460 non-null int64
Functional        1460 non-null object
Fireplaces        1460 non-null int64
FireplaceQu       770 non-null object
GarageType        1379 non-null object
GarageYrBltd      1379 non-null float64
GarageFinish      1379 non-null object
GarageCars        1460 non-null int64
GarageArea        1460 non-null int64
GarageQual        1379 non-null object
GarageCond        1379 non-null object
PavedDrive        1460 non-null object
WoodDeckSF        1460 non-null int64
OpenPorchSF       1460 non-null int64
EnclosedPorch     1460 non-null int64
3SsnPorch         1460 non-null int64
ScreenPorch       1460 non-null int64
PoolArea          1460 non-null int64
PoolQC            7 non-null object
Fence             281 non-null object
MiscFeature       54 non-null object
MiscVal           1460 non-null int64
MoSold            1460 non-null int64
YrSold            1460 non-null int64
SaleType          1460 non-null object
SaleCondition     1460 non-null object
SalePrice         1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

```

In [5]: d_missing=d.isna().sum()

missing[d_missing>0].sort_values(ascending=False)

In [6]: d_missing[d_missing>0].sort_values(ascending=False)

Out[6]: PoolQC            1453
        MiscFeature      1406
        Alley           1369
        Fence           1179
        FireplaceQu       690

```

```

LotFrontage      259
GarageYrBlt       81
GarageType        81
GarageFinish      81
GarageQual        81
GarageCond        81
BsmtFinType2      38
BsmtExposure      38
BsmtFinType1      37
BsmtCond          37
BsmtQual          37
MasVnrArea        8
MasVnrType        8
Electrical        1
dtype: int64

```

In [7]: *#keeping only columns which dont have na*

```

In [8]: d=d.dropna(axis=1, how='any')
        d.shape

```

Out [8]: (1460, 62)

```

In [9]: #removing Id field which doesnt have impact on house price.
        #del d['Id']
        d.head()

```

```

Out [9]:   Id  MSSubClass MSZoning  LotArea  Street  LotShape  LandContour  Utilities  \
0    1         60      RL      8450   Pave      Reg          Lvl      AllPub
1    2         20      RL      9600   Pave      Reg          Lvl      AllPub
2    3         60      RL     11250   Pave      IR1          Lvl      AllPub
3    4         70      RL      9550   Pave      IR1          Lvl      AllPub
4    5         60      RL     14260   Pave      IR1          Lvl      AllPub

      LotConfig  LandSlope  ...  EnclosedPorch  3SsnPorch  ScreenPorch  PoolArea  \
0      Inside      Gtl    ...              0           0           0           0
1         FR2      Gtl    ...              0           0           0           0
2      Inside      Gtl    ...              0           0           0           0
3      Corner      Gtl    ...            272           0           0           0
4         FR2      Gtl    ...              0           0           0           0

      MiscVal  MoSold  YrSold  SaleType  SaleCondition  SalePrice
0           0        2    2008        WD          Normal    208500
1           0        5    2007        WD          Normal    181500
2           0        9    2008        WD          Normal    223500
3           0        2    2006        WD          Abnorml    140000
4           0       12    2008        WD          Normal    250000


```

[5 rows x 62 columns]

```
In [10]: #correlation matrix
```

```
In [86]: import seaborn as sns
import matplotlib.pyplot as plt
matrix = d.corr()
f, ax = plt.subplots(figsize=(16, 12))
sns.heatmap(matrix, vmax=0.7, square=True)
```

```
Out[86]: <matplotlib.axes._subplots.AxesSubplot at 0x1ad3cb39e80>
```



```
In [87]: #selcting only features which are highly correlated
tcf=matrix['SalePrice'].sort_values(ascending=False)
```

```
In [88]: # Filter out the target variables (SalePrice) and variables with a low correlation score
tcf = tcf[abs(tcf) >= 0.6]
```

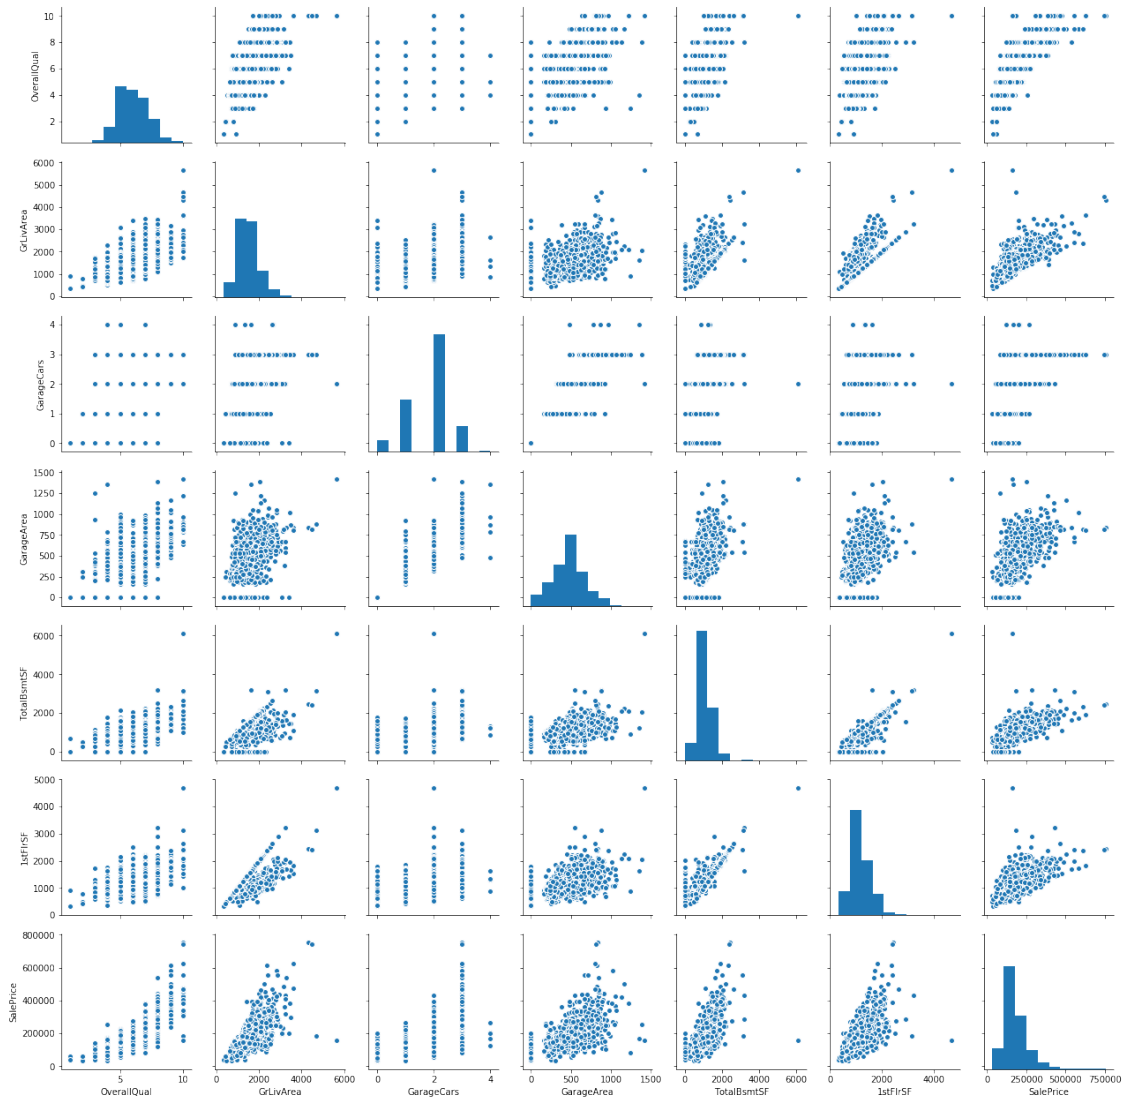
```
In [89]: tcf = tcf[tcf.index != 'SalePrice']
tcf
```

```
Out[89]: OverallQual    0.790982
         GrLivArea      0.708624
         GarageCars      0.640409
         GarageArea      0.623431
         TotalBsmtSF     0.613581
         1stFlrSF        0.605852
         Name: SalePrice, dtype: float64
```

```
In [90]: tcf.shape
```


```
Out[90]: (6,)
```

```
In [91]: cols = tcf.index.values.tolist() + ['SalePrice']
         sns.pairplot(d[cols], size=2.5)
         plt.show()
```



```
In [94]: # Build the correlation matrix
matrix = d[cols].corr()
f, ax = plt.subplots(figsize=(8, 6))
sns.heatmap(matrix, vmax=1.0, square=True)
```

```
Out[94]: <matplotlib.axes._subplots.AxesSubplot at 0x1ad3bf9e550>
```



```
In [146]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
X=d1.loc[:,d1.columns!='SalePrice']
y = d1['SalePrice']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```




```
Out[146]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                                oob_score=False, random_state=42, verbose=0, warm_start=False)
```

```
In [147]: y_pred = model.predict(X_test)
```

```
# Build a plot
plt.scatter(y_pred, y_test)
plt.xlabel('Prediction')
plt.ylabel('Real value')

# Now add the perfect prediction line
diagonal = np.linspace(0, np.max(y_test), 100)
plt.plot(diagonal, diagonal, '-r')
plt.show()
```



```
In [148]: from sklearn.metrics import mean_squared_log_error, mean_absolute_error

print('MAE:\t$%.2f' % mean_absolute_error(y_test, y_pred))
print('MSLE:\t$%.5f' % mean_squared_log_error(y_test, y_pred))
```

MAE: \$26348.38
MSLE: 0.04555

```
In [149]: #Score/Accuracy  
          print("Accuracy --> ", model.score(X_test, y_test)*100)
```

Accuracy --> 0.684931506849315

```
In [153]: #Train the model  
          from sklearn import linear_model  
          model = linear_model.LinearRegression()
```


```
In [154]: #Fit the model  
          model.fit(X_train, y_train)
```

Out[154]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

```
In [157]: #Score/Accuracy  
          print("Accuracy --> ", model.score(X_test, y_test)*100)
```

Accuracy --> 79.24553693088554

```
In [159]: # Build a plot  
          plt.scatter(y_pred, y_test)  
          plt.xlabel('Prediction')  
          plt.ylabel('Real value')  
  
          # Now add the perfect prediction line  
          diagonal = np.linspace(0, np.max(y_test), 100)  
          plt.plot(diagonal, diagonal, '-r')  
          plt.show()
```



In [150]: *#Train the model*

```
from sklearn.ensemble import GradientBoostingRegressor
GBR = GradientBoostingRegressor(n_estimators=100, max_depth=4)
```

In [151]: *#Fit*

```
GBR.fit(X_train, y_train)
```

Out[151]: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None, learning_rate=0.1, loss='ls', max_depth=4, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)


In [152]: `print("Accuracy --> ", GBR.score(X_test, y_test)*100)`

Accuracy --> 87.72567683930332

In [158]: *# Build a plot*

```
plt.scatter(y_pred, y_test)
plt.xlabel('Prediction')
plt.ylabel('Real value')
```

```
# Now add the perfect prediction line  
diagonal = np.linspace(0, np.max(y_test), 100)  
plt.plot(diagonal, diagonal, '-r')  
plt.show()
```



In []: